

06 - CE1006 - 15th Feb - Conditional Constructs

Branchless Logic (If Else)

- Avoid using conditional jumps as jump flushes the pipeline
- How?
 - Exploit Arithmetic relationship
 - AND & ROR
 - Conditional Execution
 - ARM Processor, MOVEQ

Switch Statement

- Sequential & Small Values
 - Keeping a Jump Table (JT) of all start addresses to case values
 - Move it into the Program Counter
 - MOV PC, [R2 + JT]
 - R2 = User Input
 - JT = Jump Table
- Random & Wide Values
 - Fork Algorithm - Divide and Conquer with nested ifs
 - It prevents the last case to be the worst case scenario

Loops

- Pre (While), Post (Do-while) tests
- While
 - Pre-test
 - Back CMP ... <Inverse>
 - JLE Exit
 - {S0}
 - JMP Back
 - Exit
- Do-while
 - Post-test
 - {S0}
 - CMP ... <No Inverse>
 - JGT Back
- For Loop
 - Pre-test
 - MOV [N], #0 <- Initialization
 - ... {Same as While}
 - INC [N]
 - JMP Back

07 - CE1006 - 1st Feb

Shift & Rotate Operations

Shift

- Left or Right and will be pushed out to Carry bit
- 0 or 1 will be used to fill the bit

Rotate (Cyclical shift) [ROL, ROR]

- Shifts LSB to MSB or MSB to LSB
- Bit will also be stored into Carry

Shift

- performs Multiplication (shift left) & Division (shift right)
- each shift is 2^n
 - e.g. shift left by 2. $2^2 = 4$

BCSR - (Bit Clear Status Register)

- Clears a particular Status Register flag

Signed Divide using Arithmetic Shift Right [RAR]

- Shift MSB Value to the next MSB (Keeps Sign Flag)

Rotate Counter [RCN]

- Sets the counter to Rotate [ROR, RAR]

Rotate Left with Carry (RLC)

- Push Carry into LSB

Program Control Instructions

Conditional Jump (Jcc)

- JEQ (JZ). Jump = 0
 - Jump when result = 0. * Note: PC is ahead on next instruction
- Use labels instead.

Module Programming

Subroutine Implementation (CALL and RET)

- Writing as Modules (Libraries)
- Designed and tested Independently
- Reduce Program Size
- Reuse in other projects

CALL - Subroutine call

- Similar to JMP with a Return

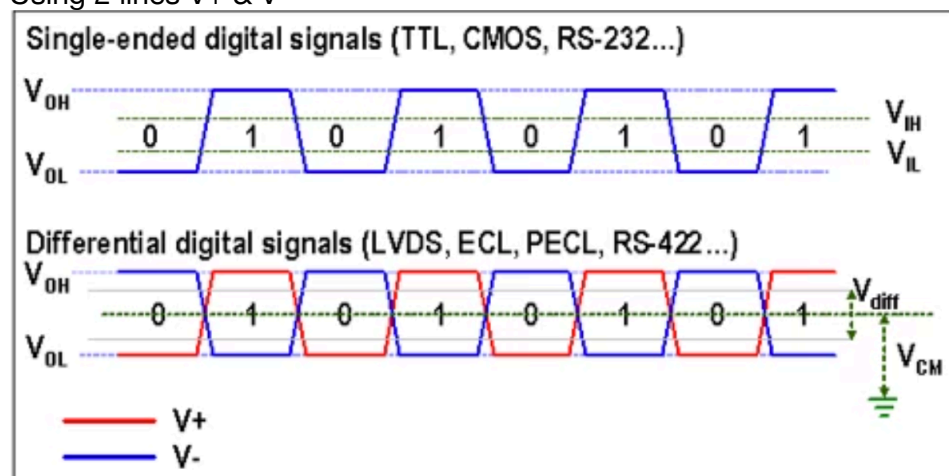
07 - CE1006 - E-Learning - 23rd Feb

Analog Signal

- Real world signals are Analog in nature
- Continuous Time and Waveform
- **Analog-to-Digital Converter (ADC):** Converts analog to digital value
- **Digital-to-Analog Converter(DAC):** converts processed signals back to Analog
- **Filter:** Smooths out DAC Signals

Interface Compatibility

- Communication Protocol
- Electrical signal level
 - Similar voltage levels between two devices
 - Transmitting Device: $V(OL)$ [0] & $V(OH)$ [1]
 - Receiving Device: $V(IL)$ [0] & $V(IH)$ [1]
 - $V(OH) \geq V(IH)$
 - $V(OL) \leq V(IL)$
 - Differential Signals
 - Better noise tolerance to allow higher clock frequency
 - Using 2 lines $V+$ & $V-$



Parallel Data Transfer

- Prone to **Signal Skew** and **Crosstalk**
- Design is simpler as Strobe Signals are needed
 - Data is being sent when Strobe Signal is asserted
- **Signal Skew**
 - data lines takes different amount of time to reach the receiver causing the wrong data to be latched
 - variation in **propagation delay**
 - contributed mainly by Capacitance and Resistance. e.g. Variations
 - $T = RC$ - Larger RC, Longer Delay
- **Crosstalk**

- Interference from neighbouring signals
- **More noise** = limits speed
- **Front-Side Bus (FSB)**
 - FSB to interface with CPU
 - has a bottleneck if connected with multiple connections

Serial Data Transfer

- **Advantage**
 - Less affected by Signal Skew & Crosstalk
- **Disadvantage**
 - Interface design is more complex as it need to handle serial to parallel conversion
- **Data Transfer**
 - Simplex - Single direction
 - Half Duplex - Both direction but only one active
 - Full Duplex - Simultaneous RX & TX
- **Synchronous vs Asynchronous**
 - Synchronous
 - Common clock signal
 - Master-Slave config: Master providing Clk
 - Potential faster transfer as no data overhead to synchronous data
 - **Inter-Integrated Circuit (I2C) Bus**
 - Half Duplex
 - Two wire bus (Multi-Master, Multi-Slave)
 - SDA: Data
 - SCL: Clock
 - Slave & Master can transmit
 - Master sets the Clock
 - Start Sequence
 - SCL: HIGH
 - SDA: Pulled LOW
 - End Sequence
 - SCL: HIGH
 - SDA: Pulled HIGH
 - Data is changed when SCL is LOW
 - Data is latched at rising edge of SCL
 - **Serial Peripheral Interface (SPI) Bus**
 - Full Duplex
 - 1 SPI Master, Multiple Slaves
 - Multiple Slave Selects on Master
 - Data Transfer: Slave Select Pulled LOW
 - MOSI: Master-Out(put)-Slave-In(put)
 - MISO: Master-In(put)-Slave-Out(put)

Asynchronous Transfer

- No common clk
- Special SYNC to indicate START/STOP

- Potential Skew as it progresses
 - Due to differences in clock
 - SYNC words/bits provide timestamp to resynchronize
- Universal Asynchronous Receiver Transmitter (UART)
 - RS232

| LOGIC LEVELS | | | |
|--------------|--------|--------------------|--------------------|
| STANDARD | | LOGIC 0 (Volts) | LOGIC 1 (Volts) |
| TTL | INPUT | 0.0 to 0.8 | 2.0 to 5.0 |
| | OUTPUT | 0.0 to 0.4 | 2.8 to 5.0 |
| RS-232 | INPUT | +3 to +15 | -3 to -15 |
| | OUTPUT | +5 to +15 | -5 to -15 |

-
- UART Transmit
 - No Clk involved
 - START Pattern (login '0') - Falling Edge
 - DATA sent out at a certain **Baud Rate**
 - STOP pattern (1 or more bits) terminates transmission
- UART Receive
 - Falling Edge trigger interrupt for data transfer
 - Receiver Clk usually run must faster.
 - e.g. 16x of Baud Rate of Transmitter to time the sampling close to the middle of each data bit
 - Oversampling increases accuracy
 - Taking more samples close to the middle of each data bit
- Universal Serial Bus (USB)
 - Grounded Shielding, Twisted pair reduce effect of Crosstalk
 - NRZI (Non-Return-to-Zero Inverted)
 - Bit Stuffing
 - Full Duplex in USB 3.0

Summary

- Interfacing Consideration
 - Electrical Signal level
 - Communication Protocol
- Types of Signal
 - Analog
 - Digital
- Digital Interface
 - Parallel
 - Serial
- Interface Standards
 - INTEL FSB
 - IEEE 488 (GPIB)
 - UART
 - USB

08 - CE1006 - 7th Mar & 10th Mar

Data Transfer

- Polling Technique - Programmed I/O
 - CPU polls certain IO port
 - CPU has full controlled dedicate 100% resources
 - Watching a bowl of water bowl
 - Advantages
 - Minimum hardware needed, easy to debug
 - Programmer has complete Control
 - Disadvantages
 - CPU stays in a loop, other tasks don't get done
 - Program execution held up waiting for device to be ready
- Interrupt
 - Signal notifies CPU when event occurs
 - CPU suspends its current program to service the interrupt
 - Interrupt Service Routine (ISR)
 - A short program to process the interrupt
 - Waiting to hear the whistle of water boiling
 - Advantages
 - Efficient use of CPU as it doesn't need to monitor I/O devices
 - Allows **prioritisation** and **pre-emption**
 - Disadvantages
 - More hardware interface which is complex and difficult to debug
- DMA Controller
 - Relieve CPU of data transfer task
 - Config
 - Src Addr
 - Dest Addr
 - Amt of data to transfer
 - DMA Trigger signal
 - Allow DMA to take over system bus
 - Design
 - Fetch-n-Deposit
 - Internal buffer to manage data transfer
 - FlyBy
 - Data transferred directly from src to dest
 - Additional Address line in DMAC
 - Modes of Transfer
 - Burst
 - DMA takes control
 - returns CPU control until transfer is completed
 - Used for high data transfer
 - Cycle Stealing
 - Releases data bus after one word transfer
 - Executed between
 - CPU Instructions
 - Pipeline stages
 - Used for responsive (real-time) systems
 - Transparent

- DMAC transfer when system bus is not in use by CPU
- CPU is not affected

10th Mar Computer Memory


- Semiconductor Based
 - SSD, SD Card
- Magnetic Based
 - Magnetic
- Optical Based
 - CD, DVD

Volatile & Non-volatile

- Volatile
 - Data lost when power is removed
 - SRAM & DRAM
- Non-volatile
 - Data retained when power is removed
 - Magnetic storage


RAM

- Static Ram
 - Data stored as long as supply is applied
 - Access
 - Addr
 - Data - Data bus
 - CS - Chip Select
 - WE - Write
 - OE - Read
 - Access via
 - Word lines
 - Horizontal
 - Controls 8 bits
 - Bit lines
 - Activates the line onto the data bus
 - SRAM Cell



- Dynamic RAM
 - Periodic refresh required
 - Stored charge leaks

- Uses a Capacitor
- Read process destroys information
- DRAM Cell



-
- Currently using Synchronous SDRAM
 - Increase speed using synchronous clock

EPROM & EEPROM

- EPROM
 - Erasable Programmable ROM
 - Put under UV light to erase program
- EEPROM
 - Electrically Erasable Programmable ROM
 - Electrically program or erase
- MOSFET
 - Works like a switch
- Floating Gate Transistor
 - Similar to MOSFET
 - Extra layer (Floating Gate)
 - Transistor Off when programmed (contains charges)
 - Logic 0
 - Transistor On when Erased (no charges)
 - Logic 1

NAND & NOR Flash

- Like EEPROM
- Erased in larger block sizes
- Erase cycle is slower than Read/Write
- NOR Flash
 - Support Execute in place
- NAND Flash
 - Does not support execute in place

09 - CE1006 - 11th Feb - Modular Programming

Calling Subroutine

- Pushing Parameter into Stack
- Using ADD SP, #2 (Pops two parameter pushed into the stack)

Transparent Subroutines

- Saving away all registers before subroutine for it to prevent overwriting registers in main program
- PSHM 0x00C (1100) <- (R3,R2,R1,R0)
- POPM to restore registers

Accessing Stack Parameters

- Finding offset from Pushing to Calling using Stack Pointer (SP)
- Note: CALL puts Return Program Counter (PC) Value into the stack
- PSHM pushes register values to save into the stack

Passing by Value or Reference

- Reference: Arrays, Modifying variable in subroutines
- Value: PSH #5
- Reference: PSH #0x100

Local Variables

- Loose coupling (Local variables) - Data within module is entirely independent of other modules
- Stack Frame: temporary memory space
 - created right after Return Address
 - Can be accessed using **Frame Pointer (FP)** - Start SP Address of subroutine

Using the Frame Pointer

- Must R3 be saved to use as Frame Pointer?
- Stack Pointer can be used
 - Pro: Efficient without the setup of Frame Pointer
 - Con: More restrictive

Summary

- CALL, RET: Basic instruct for subroutines
- Stack: multiple parameter passing
- Transparent Subroutines with PSHM, POPM
- Local Variables in subroutine using a **Stack Frame**

Control Flow Constructs

If Statements

- In more than condition, find the inverse for a more efficient code
- Using less than and equal
 - CMP a, b
 - JLE skip <- Jump if less than equal
 - {If code}
- If-Else
 - Jump to the Else situation
 - Note: **JMP required** when {If code} is completed
 - JNE DoElse
 - {If Code}
 - JMP skip
 - Do Else {Else Code}


Compounded Conditions

- AND (&&) Conditions
 - C Code: if ((a == b) && (b > 0)) {S1}
 - Low Level:
 - if (a != b) then Skip
 - if (b <= 0) then Skip
 - {S1}
 - Skip ...
 - Note: Keeping the least likely condition leftmost makes the program more efficient
- OR (||) Conditions
 - C Code: if ((a == b) || (b > 0)) {S1}
 - Low Level:
 - if (a == b) then Dolf
 - if (b <= 0) then Skip <- **Reversed Condition**
 - Dolf {S1}
 - Skip ...

09 - CE1006 - Magnetic Hard Disk

Types of Recording

- Longitudinal Recording
- Perpendicular Recording



Dissection of Magnetic HDD

- Platters
- Tracks
- Sectors


HDD Transfer Rate

- Seek Time T_S
 - Time Taken for head to move to correct track
- Rotational delay T_R
 - Time taken for disk to rotate till start of target sector
 - Revolution Per Minute (RPM) usually converted to RPS
 - $RPM / 60 = RPS$
 - $T_R = 0.5 / RPS$ seconds
- Access Time T_A
 - Time from request till head in position
 - $T_A = T_S + T_R$
- Transfer Time T_T
 - Time to transfer required data after head is positioned
 - $T_T = N / RPS * D_T * D_S$
 - N = Number of bytes to read
 - D_T = Number of sectors per track
 - D_S = number of bytes per sector
- Total Time

- $T_{\text{Total}} = T_A + T_T$
- It is notable that Access time is much higher than Transfer time thus the slow speed when hard disk is fragmented

Physical Layout - Hardware View

- Tracks divided into zone, with different number of sectors / track



Logical Layout - Software View

- Cylinder-Head-Sector (CHS)
 - Legacy Scheme due to unable to meet current space standards
- Logical Block Addressing (LBA)
 - 48 bit allowing up to 128 PetaByte

Datasheet

- Complexity of Hard disk is expanding
- Most useful resource are
 - Bytes per Sector
 - Number of Sectors
 - Total Logical Data Bytes

Solid State Drive (SSD)

- Based on NAND-Flash or NOR-Flash
- Limited Program-erase cycles (3,000 - 1,000,000)
- Extending Techniques
 - Wear Leveling
 - External RAM as buffer
- Mean time between failures (MTBF) for SSD is comparable to HDD


HDD vs SSD

- HDD
 - Advantages
 - Lower Cost per Bit
 - Almost infinite erase cycle
 - Cons
 - more prone to crashing if HDD is dropped or shaken
 - Slower Transfer Rate
- SSD

- Advantages
 - Lighter and occupy less space
 - High Transfer Rate
- Disadvantages
 - More Costly compared to HDD
 - Finite Erase Cycles

Reliability – Bath Tub Curve

- Initial Failure
- Random Failure
- Wear-out Failure



- Burn in - Stress a product
- MTTF - Mean (Average) Time To Failure


$$R(t) = e^{-\frac{\text{time}}{MTTF}}$$

$$R(t) = e^{-\frac{MTTF}{MTTF}} = e^{-1} = 0.368$$

Redundancy and RAID

- Redundant Array of Independent Disks
- RAID 0
 - Performance focused
 - No redundancy
- RAID 1
 - Mirrors Hard disk
 - Faster Read Operations
 - Write throughput is slower
- RAID 5
 - Higher performance
 - Stripping of file with distributed parity (Error Correction)
 - RAID 3 - 4 Similar to 5

- RAID 6
 - Recover from multiple disk failures
 - Double distributed parity
- RAID 10 (Not Official)
 - Combination of RAID 1 + RAID 0



◦


09 - CE1006 - System Memory

Cache

- Buffer between CPU & Main Memory
- Ghz Speed
- External Memory 100s Mhz
- speed up recently used data
- Common Data Mapping Scheme

CPU Memory Access

- CPU sends request to Cache
- If not in cache, request from Main Memory, if not from Disk
- Once located, required data and **nearby data elements** are fetched into cache
- whole cache block is transferred when one or more bytes required



Common Data Mapping Scheme

- **Direct Mapped Cache**
 - One block can only go into one cache block location
 - Corresponding **Tag Value** keeps track of the main memory block
 - Tag, Block, Offset
 - MM Address
 - Advantages
 - Simple
 - Found easily and rapidly
 - Disadvantages
 - Prone to Cache Trashing
- **N-way Set Associative Cache**
 - e.g. **2-Way** Set Associative Cache
 - N Blocks forms a Set, 2 Blocks forms a set
 - Tag, Set, Offset
 - Advantages
 - Less Cache Trashing

- Disadvantage
 - Longer search time
 - Complex Design
- **Fully Associative Cache**
 - Data can be in any Cache Block
 - Tag, Offset
 - Extreme case of N-Way Set Associative Cache

Principles of Cache Design

- **Storage**
 - Principle of locality
 - Due to array usage, chances of using a **nearby data element** accessed is high
 - Locality of Space (Spatial Locality)
 - Code/Data close together
 - Locality of Time (Temporal Locality)
 - Recently executed code
 - Cache Memory Replacement Policy
 - Least Recently Used (LRU) Algorithm
 - First-in, First-out (FIFO)
 - Random
 - Cache Write Policy
 - Dirty Blocks
 - Bytes changed in cache
 - Write Through
 - updates cache and main memory simultaneously
 - Write Back
 - Updates memory only when block is to be replaced
 - Write Miss Policy
 - Write Allocate??
 - Write-no-allocate
 - Write directly to main memory
- **Retrieval**
 - Cache Hit - Data in Cache
 - Cache Miss - Data not in Cache
 - Cache Hit Rate - Percentage of time data in cache
 - Cache Miss Rate - Percentage of time data not in cache
 - Miss Rate = 1 - Hit Rate
 - Hit Time - Time to access cache memory
 - Miss penalty - Time to process a miss
 - Effective Access Time
 - $EAT = H \times Access_c + (1-H) \times Miss\ Penalty$
 - $EAT = H \times Access_c + (1-H) \times (Access_c + Access_{MM})$
 -


09 - CE1006 - Virtual Memory Management

Virtual Memory

- Allows programs bigger than size of main memory to be processed
- Memory Management Unit (MMU)
 - Hardware Device that does conversion
- Paging
- Segmentation
- Advantages
 - Frees application from managing shared memory space
 - Virtual Address Space to appear contiguous on fragmented blocks

Paging - Not Understood

- Page Table to Convert
- Frame - Physical Page numbers in main memory
- Page - Virtual Page numbers in program code
- Page Fault - Data not found in main memory
- Reference to only Valid Bit 1
- Offsets remain the SAME




- Translation Lookaside Buffer (TLB)
 - Page Table Cache that stores recently accessed values

Segmentation Method

- Variable-length segments
- entries in a **Segment Table**
- Note: Still **Page Fault**, not segmentation fault

Memory Fragmentation

- Memory too small to contain any data and rendered useless
- Paging - Internal Fragmentation
- Segmentation - External Fragmentation



11 - CE1006 - 18th Feb - Data Transfer and I/O Interfaces

Lab 3 - UART Communication

Interface

- Boundary where two or more devices exchange data
 - Networks: Peer-To-Peer or Multi-point
 - Data Transfer: Serial(Single Bit) or Parallel(Multiple Bits)
 - Single or Bi Direction

Serial Interface

- Transfer Modes
 - Simplex - One Sender, Receiver
 - Half Duplex - One Line, Sender/Receiver
 - e.g. Walkie Talkie
 - Full Duplex - Two Lines for TX & RX
 - e.g. - Handphone

Asynchronous vs Synchronous

- Asynchronous
 - No Common Clock
 - Receiver needs to know transmitting clock rate
 - Special SYNC words to track START/STOP
- UART - Universal asynchronous receiver/transmitter
 - Baud Rate - Transmitting clock rate
 - START pattern
 - PARITY bit (optional)
 - STOP pattern

Parity Bit

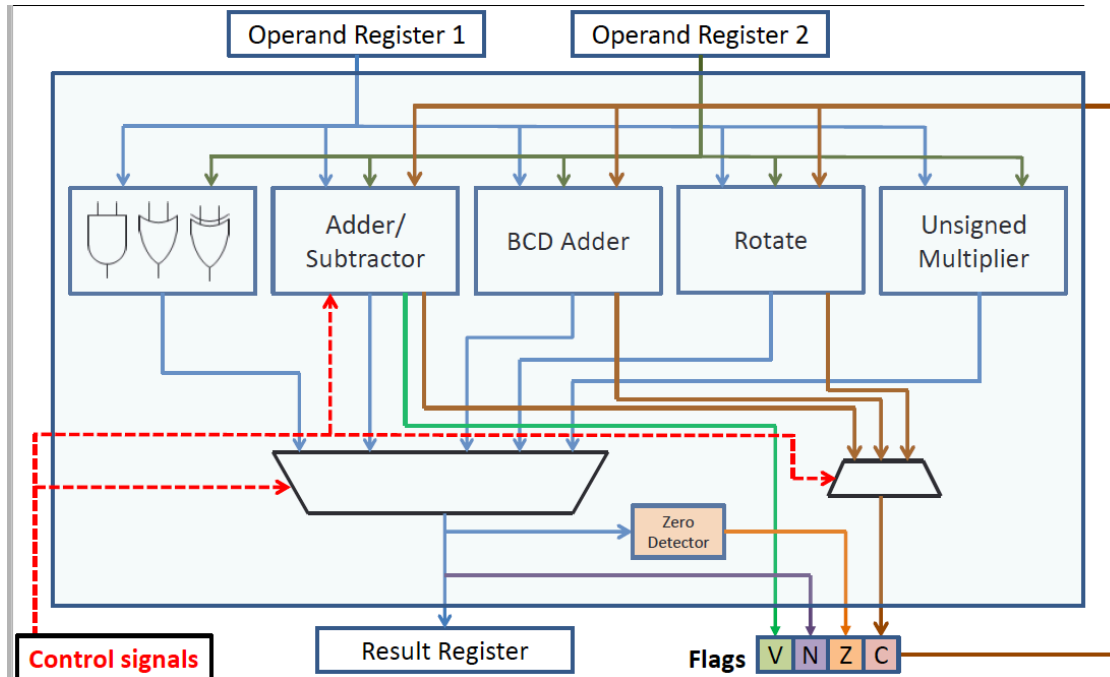
- Even Parity
 - 0101 - Parity Bit is one to make all even 1 bits

11 - CE1006 - Arithmetic and Logic Unit

Arithmetic and Logic Unit

- Arithmetic - Addition, Subtraction
- Logic - NOT, OR ...

- ALU of VIP Processor



Positional Numbering System

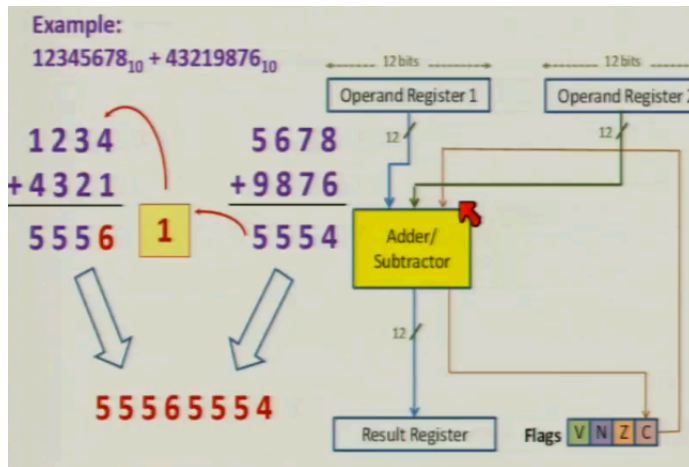
- As number gets larger, the benefits increases

Decimal to Binary Conversion by Subtraction

Signed Magnitude vs 2's complement

- Unsigned Magnitude
 - Overflow when carry is 1
- Signed Magnitude
 - 0111 = 8
 - 1111 = -8
 - Easy for people to understand
 - Requires complicated hardware
 - **Two representations of zero**
- Two's Complement
 - Detecting Overflow
 - Same sign and result sign in inverted sign
 - Carry does not determine overflow


Multi-Precision Arithmetic



Binary Multiplication

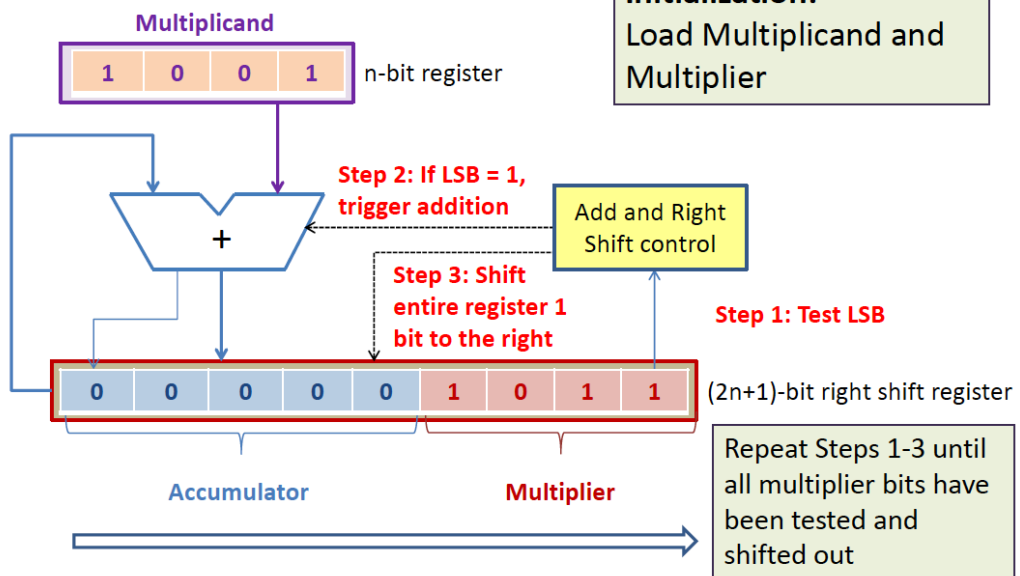
- Successive Addition
 - $4 * 3 = 4 + 4 + 4$
- Partial Product

Example: $9_{10} \times 11_{10}$ (unsigned)



- Multiplier, Multiplicand, Result, Number of Multiplier bits
- UMUL <Multiplier> in VIP
- Multiplicand in R0
- Hardware Multiplier

Example: $9_{10} \times 11_{10}$ (unsigned)



-
- **Booth's Algorithm**
 - Optimise number of addition cases

| Current Bit | Preceding Bit | Booth's Form |
|-------------|---------------|--------------|
| 1 | 0 | -1 |
| 0 | 1 | +1 |
| 0 | 0 | 0 |
| 1 | 1 | 0 |

- Add virtual 0 behind LSB

Understand Signed Multiplication
Multiply Negative Numbers

11 - CE1006 - Floating Point Unit

Range

- Smallest and largest representable number
- Each tick represents a number in the range

Precision

- Amount of information to represent each number
- Number of ticks represent the precision

Radix Point

- dot to represent the point between Range and Precision

Number of bits are the same

- $\text{Precision} = 1 / \text{Range}$
- Range increase, Precision Decreases

Floating Point Representation

- Sign
 - Sign bit
- Mantissa
 - Base Value
 - Size determine precision of values
- Exponent
 - Specifies decimal of radix point
 - Size determine range of representable numbers
- Small Numbers require high precision, Large numbers require lesser precision

Normalisation

- Single integer cannot be zero

Underflow

- When number is too small to be represented

Guard Bit and Rounding

- Guard Bit
 - When add/sub exponents are required to be aligned and some bits might be lost
 - Maintain precision during computation
- Rounding

- Remove the Guard Bit by rounding
- Rounding Error
 - Results are inaccurate due to limited bits
 - Adding a list of very small numbers
 - Workarounds
 - Add/Sub similar exponent
 - Start with smaller exponents

Multiplication

- Multiply Mantassa
- Add Exponents
- Normalise if required

IEEE 754 Floating Point Standard

- Provides high precision with large range
- Modes

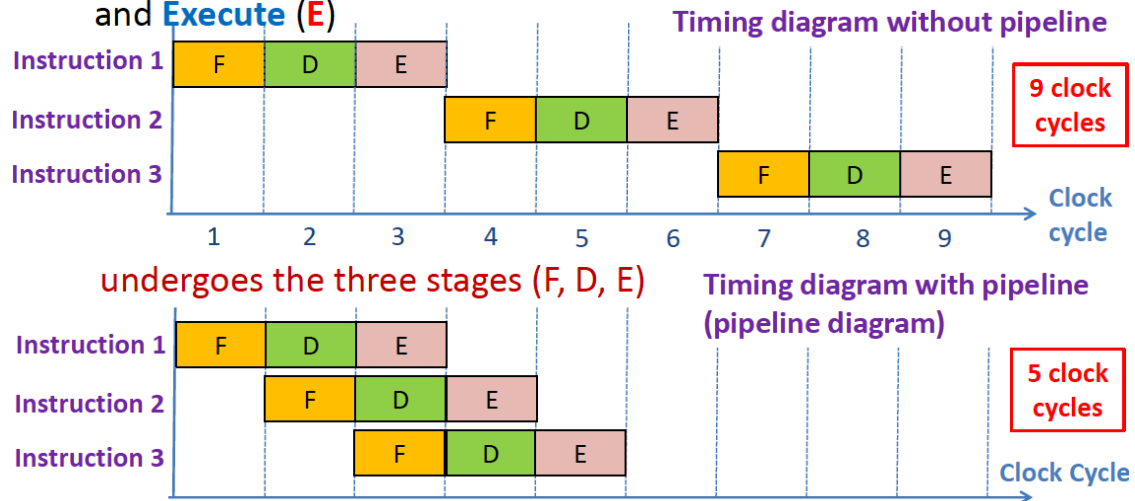
| Mode | Sign | Exponent | Fraction |
|--------------------|-------|----------------------|---------------|
| Normalized | 1 / 0 | 00000001 to 11111110 | Anything |
| Denormalized | 1 / 0 | 00000000 | Non zero |
| Zero | 1 / 0 | 00000000 | 0000 ... 0000 |
| Infinity | 1 / 0 | 11111111 | 0000 ... 0000 |
| Not a Number (NaN) | 1 / 0 | 11111111 | Non zero |

-
- Components
 - Sign
 - Exponent
 - Bias = 127 (Single), 1023 (Double)
 - Exponent = 101100102 = 178; E – Bias = 178 – 127 = 51
 - Fraction
 - Assumes hidden 1
 - 0110 = 1.0110
- Single Precision
 - 32 bits
 - 1,8,23
- Double Precision
 - 64 bits
- Different modes to fix underflow
 - Normalized
 - De Normalized
- Conversion
 - Fraction x $2^{\text{Exponent} - \text{Bias}}$

12 - CE1006 - Parallelism

CPU

- Instructions are in 3 stages
 - Fetch
 - Decode
 - Execute
- By splitting up the resource, more instructions can be executed in parallel and **Execute (E)**




Pipelining Conflicts

- Minimum Time
 - No Pipeline Conflicts that slows execution
- Conditions
 - Resource conflict
 - Two instructions requesting same resource at the same time
 - Resolutions
 - Wait
 - Hold instruction until delay is over
 - Requires mechanism to check for conflict
 - Add Resources
 - Add more hardware
 - Data dependencies
 - Instruction depends on previous result
 - Resolution
 - Stall pipeline - Hardware (Circuitry)
 - Insert NOP instructions - Software (Compiler)
 - Branch statements
 - Evaluate condition to know which branch to be taken

- Unnecessary instructions are fetched if branch jumps
- Resolutions

- Reducing Branch Delay
 - Calculating Branch target at Decode stage
 - Requires extra adder in circuitry
 - Reduces branch delay to 1 cycle



- Delayed Branching
 - Fetch the instruction directly before the branch
 - Ensure no branch delay
 - Note: Instruction must be independent and does not affect the branching value
 - Compiler can schedule such instructions
 - If no independent instructions
 - one cycle delay
 - NOP instruction added
- Dynamic Branch Prediction
 - Hardware circuit to predict outcome
 - Prediction
 - Correct - No wasted cycles
 - Incorrect - Wasted cycles
 - Update prediction bit for future use